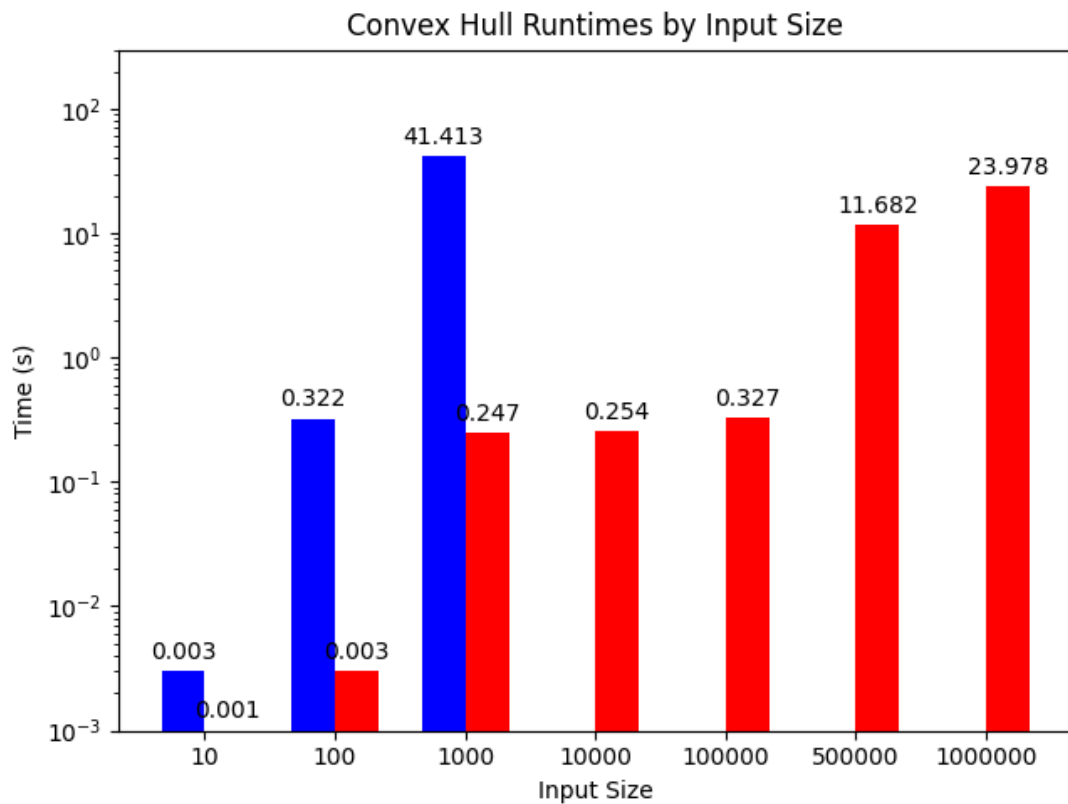


Runtime Analysis		
Input Size	Brute Force Runtime	Graham Scan Runtime
n = 10	00:00.003	00:00.001
n = 100	00:00.322	00:00.003
n = 1,000	00:41.413	00:00.247
n = 10,000	> 5 minutes	00:00.254
n = 100,000	> 5 minutes	00:02.327
n = 500,000	> 5 minutes	00:11.682
n = 1,000,000	> 5 minutes	00:23.978



Brute force runtimes are shown in blue, and Graham Scan runtimes are shown in red. Runtimes over 5 minutes were not plotted. The y-axis is shown with a logarithmic scale so that all runtimes are visible.

## Methodology & Analysis

Runtimes listed are the average from running the program three times with the same input sets. Points in each input set were randomly generated. I stopped computation early if the program took more than 5 minutes to run. To generate 500,000 and 1,000,000 points, I had to remove the logic disallowing duplicates because it took too long.

The brute force algorithm took longer than 5 minutes when I tried to run it with an input size of  $n = 10,000$ . In contrast, the Graham Scan method took less than 30 seconds, even when testing with 1,000,000 points. This clearly illustrates the importance of an efficient algorithm and the difference between  $O(n^3)$  and  $O(n \log n)$  efficiency, as well as how quickly a runtime spirals out of control with an  $O(n^3)$  runtime.